Frequently asked questions

Last Updated: 2023-08-31

Find answers to frequently asked questions about watsonx.ai.

IBM watsonx.ai questions 🔗

- How do I sign up for watsonx.ai?
- Can I try watsonx.ai for free?
- How do I upgrade watsonx.ai?
- Which regions can I provision watsonx.ai in?
- Which web browsers are supported for watsonx.ai?
- How can I get the most runtime from my Watson Studio Lite plan?
- How do I change languages for the product and the documentation?
- How do I find my IBM Cloud account owner or administrator?
- Can I provide feedback?

Foundation model questions 🔗

- What foundation models are available and where do they come from?
- What data was used to train foundation models?
- Do I need to check generated output for biased, inappropriate, or incorrect content?
- Is there a limit to how much text generation I can do?
- Does prompt engineering train the foundation model?
- Why can't I find the Prompt Lab?
- Does IBM have access to or use my data in any way?
- What APIs are available?

Project questions 🔗

- How do I load very large files to my project?

IBM Cloud Object Storage questions 🔗

- What is saved in IBM Cloud Object Storage for workspaces?
- Do I need to upgrade IBM Cloud Object Storage when I upgrade other services?

Notebook questions ⌘

Security and reliability questions ⌘

Sharing and collaboration questions ⌘

Machine learning questions ⌘

IBM watsonx.ai questions ⌘

How do I sign up for watsonx.ai? ⌘

Go to Try IBM watsonx.ai.

Can I try watsonx.ai for free? ⌘

Yes, when you sign up for IBM watsonx.ai, you automatically provision the free version of the underlying services: Watson Studio and Watson Machine Learning.

How do I upgrade watsonx.ai? 🔗

When you're ready to upgrade any of the underlying services for watsonx.ai, you can upgrade in place without losing any of your work or data.

You must be the owner or administrator of the IBM Cloud account for a service to upgrade it. See Upgrading services on watsonx.

How can I get the most runtime from my Watson Studio Lite plan? 🔗

The Watson Studio Lite plan allows for 10 CUH per month. You can maximize your available CUH by setting your assets to use environments with lower CUH rates. For example, you can change your notebook environment. To see the available environments and the required CUH, go to the Services catalog page for Watson Studio.

How do I find my IBM Cloud account owner? 🔗

If you have an enterprise account or work in an IBM Cloud that you don't own, you might need to ask an account owner to give you access to a workspace or another role.

To find your IBM Cloud account owner:

1. From the navigation menu, choose **Administration > Access (IAM)**.
2. From the avatar menu, make sure you're in the right account, or switch accounts, if necessary.
3. Click **Users**, and find the username with the word owner next to it.

To understand roles, see Roles for IBM watsonx. To determine your roles, see Determine your roles.

Can I provide feedback? 🔗

Yes, we encourage feedback as we continue to develop this platform. From the navigation menu, select **Support > Share an idea**.

Foundation models 🔗

What foundation models are available and where do they come from? 🔗

See the complete list of supported foundation models.

What data was used to train foundation models? ⊘

Links to details about each model, including pretraining data and fine-tuning, are available here: Supported foundation models.

Do I need to check generated output for biased, inappropriate, or incorrect content? ⊘

Yes, you must review the generated output of foundation models. Third Party models have been trained with data that might contain biases and inaccuracies and can generate outputs containing misinformation, obscene or offensive language, or discriminatory content.

In the Prompt Lab, when you toggle **AI guardrails** on, any sentence in the prompt text or model output that contains harmful language will be replaced with a message saying potentially harmful text has been removed.

See Avoiding undesirable output.

Is there a limit to how much text generation I can do? ⊘

With the free trial of watsonx.ai, you can use up to 25,000 tokens per month. Your token usage is the sum of your input and output tokens.

With a paid service plan, there is no token limit, but you are charged for the tokens that you submit as input plus the tokens that you receive in the generated output.

See Watson Machine Learning plans.

Does prompt engineering train the foundation model? ⊘

No, submitting prompts to a foundation model does not train the model. The models available in watsonx.ai are pretrained, so you do not need to train the models before you use them.

See Prompt tips.

Why can't I find the Prompt Lab? ⊘

The Prompt Lab is currently available only in the Dallas and Frankfurt service regions. If you provisioned watsonx.ai in another region, you won't see the Prompt Lab option on the **New task** page within a project. You must sign up for watsonx.ai again, but this time in the Dallas or Frankfurt regions. Go to the appropriate URL for your current region to sign up in Dallas or Frankfurt:

– London region

– [Tokyo region](#)

Because you can have only one instance of a Lite service, your existing instances of Watson Studio and Watson Machine Learning for your current region are deleted.

If you provisioned watsonx.ai in another region, and you click the **Experiment with foundation models and build prompts** tile on the watsonx.ai home page, you are asked to sign up for watsonx.ai again in the Dallas or Frankfurt regions.

Does IBM have access to or use my data in any way? ✎

No, IBM does not have access to your data.

Your work on watsonx.ai, including your data and the models that you create, are private to your account:

– Your data is accessible only by you. Your data is used to train only your models. Your data will never be accessible or used by IBM or any other person or organization. Your data is stored in dedicated storage buckets and is encrypted at rest and in motion.
– Your models are accessible only by you. Your models will never be accessible or used by IBM or any other person or organization. Your models are secured in the same way as your data.

Learn more about security and your options:

– [Security and privacy of foundation models](#)
– [Security for IBM watsonx.ai](#)
– [Data security](#)

What APIs are available? ✎

You can prompt foundation models in watsonx.ai programmatically using the Python library.

See [Foundation models Python library](#).

Projects ✎

How do I load very large files to my project? ✎

You can't load data files larger than 5 GB to your project. If your files are larger, you must use the Cloud Object Storage API and load the data in multiple parts. See the [curl commands](#) for working with Cloud Object Storage directly on IBM Cloud.

See [Adding very large objects to a project's Cloud Object Storage](#).

IBM Cloud Object Storage ⊘

What is saved in IBM Cloud Object Storage for workspaces? ⊘

When you create a project or deployment space, you specify a IBM Cloud Object Storage and create a bucket that is dedicated to that workspace. These types of objects are stored in the IBM Cloud Object Storage bucket for the workspace:

– Files for data assets that you uploaded into the workspace.

– Files associated with assets that run in tools, such as, notebooks and models.

– Metadata about assets, such as the asset type, format, and tags.

Do I need to upgrade IBM Cloud Object Storage when I upgrade other services? ⊘

You must upgrade your IBM Cloud Object Storage instance only when you run out of storage space. Other services can use any IBM Cloud Object Storage plan and you can upgrade any service or your IBM Cloud Object Storage service independently.

Why am I unable to add storage to an existing project or to see the IBM Cloud Object Storage selection in the New Project dialog? ⊘

IBM Cloud Object Storage requires an extra step for users who do not have administrative privileges for it. The account administrator must enable nonadministrative users to create projects.

If you have administrator privileges and do not see the latest IBM Cloud Object Storage, try again later because server-side caching might cause a delay in rendering the latest values.

Notebooks ⊘

Can I install libraries or packages to use in my notebooks? ⊘

You can install Python libraries and R packages through a notebook, and those libraries and packages will be available to all your notebooks that use the same environment template. For instructions, see Import custom or third-party libraries. If you get an error about missing operating system dependencies when you install a library or package, notify IBM Support. To see the

preinstalled libraries and packages and the libraries and packages that you installed, from within a notebook, run the appropriate command:

– **Python**: `!pip list`

– **R**: `installed.packages()`

Can I call functions that are defined in one notebook from another notebook? ∅

There is no way to call one notebook from another notebook on the platform. However, you can put your common code into a library outside of the platform and then install it.

Can I add arbitrary notebook extensions? ∅

No, you can't extend your notebook capabilities by adding arbitrary extensions as a customization because all notebook extensions must be preinstalled.

How do I access the data from a CSV file in a notebook? ∅

After you load a CSV file into object storage, load the data by clicking the **Code snippets** icon (`</>`) in an opened notebook, clicking **Read data** and selecting the CSV file from the project. Then, click in an empty code cell in your notebook and insert the generated code.

How do I access the data from a compressed file in a notebook? ∅

After you load the compressed file to object storage, get the file credentials by clicking the **Code snippets** icon (`</>`) in an opened notebook, clicking **Read data** and selecting the compressed file from the project. Then, click in an empty code cell in your notebook and load the credentials to the cell. Alternatively, click to copy the credentials to the clipboard and paste them into your notebook.

Security and reliability ∅

How secure is IBM watsonx? ∅

The IBM watsonx platform is very secure and resilient. See [Security of IBM watsonx](#).

Is my data and notebook protected from sharing outside of my collaborators? ∅

The data that is loaded into your project and notebooks is secure. Only the collaborators in your project can access your data or notebooks. Each platform account acts as a separate tenant of the Spark and IBM Cloud Object Storage services. Tenants cannot access other tenant's data.

If you want to share your notebook with the public, then hide your data service credentials in your notebook. For the Python and R languages, enter the following syntax: `# @hidden_cell`

Be sure to save your notebook immediately after you enter the syntax to hide cells with sensitive data.

Only then should you share your work.

Do I need to back up my notebooks? ⅋

No. Your notebooks are stored in IBM Cloud Object Storage, which provides resiliency against outages.

Sharing and collaboration ⅋

What are the implications of sharing a notebook? ⅋

When you share a notebook, the permalink never changes. Any person with the link can view your notebook. You can stop sharing the notebook by clearing the checkbox to share it. Updates are not automatically shared. When you update your notebook, you can sync the shared notebook by reselecting the checkbox to share it.

How can I share my work outside of RStudio? ⅋

One way of sharing your work outside of RStudio is connecting it to a shared GitHub repository that you and your collaborators can work from. Read this blog post for more information.

However, the best method to share your work with the members of a project is to use notebooks in the project that uses the R kernel.

RStudio is a great environment to work in for prototyping and working individually on R projects, but it is not yet integrated with projects.

How do I share my SPSS Modeler flow with another project? ⅋

By design, modeler flows can be used only in the project where the flow is created or imported. If you need to use a modeler flow in a different project, you must download the flow from current project (source project) to your local environment and then import the flow to another project (target project).

IBM Watson Machine Learning 🔗

How do I run an AutoAI experiment? 🔗

Go to Creating an AutoAI experiment from sample data to watch a short video to see how to create and run an AutoAI experiment and then follow a tutorial to set up your own sample.

What is available for automated model building? 🔗

The AutoAI graphical tool automatically analyzes your data and generates candidate model pipelines that are customized for your predictive modeling problem. These model pipelines are created iteratively as AutoAI analyzes your data set and discovers data transformations, algorithms, and parameter settings that work best for your problem setting. Results are displayed on a leaderboard, showing the automatically generated model pipelines ranked according to your problem optimization objective. For details, see AutoAI overview.

What frameworks and libraries are supported for my machine learning models? 🔗

You can use popular tools, libraries, and frameworks to train and deploy machine learning models by using IBM Watson Machine Learning. The supported frameworks topic lists supported versions and features, as well as deprecated versions scheduled to be discontinued.

What is an API Key? 🔗

API keys allow you to easily authenticate when using the CLI or APIs that can be used across multiple services. API Keys are considered confidential since they are used to grant access. Treat all API keys as you would a password since anyone with your API key can impersonate your service.